

# Phase 1: Infrastructure Setup

This phase covers the creation of the necessary cloud servers.

## 1. Launch Jenkins Master & Build Slave:

- Launch two AWS EC2 instances (type `t2.micro` is sufficient).
- Name them `Jenkins-Master` and `Build-Slave`.
- Use an **Amazon Linux 2** AMI.
- Ensure the security group allows SSH (port 22) access.

## 2. Install Jenkins on Master:

- SSH into the `Jenkins-Master`.
- Install Java 17: `sudo yum install java-17-amazon-corretto -y`
- Install Jenkins and start the service as we did in the first project.
- Complete the initial Jenkins setup wizard.

## 3. Prepare Build Slave:

- SSH into the `Build-Slave`.
  - Install Java 17: `sudo yum install java-17-amazon-corretto -y`
  - Install Docker: `sudo yum install docker -y`
  - Start and enable Docker: `sudo systemctl start docker && sudo systemctl enable docker`
  - Add user to Docker group: `sudo usermod -aG docker ec2-user` (then log out and log back in).
  - Connect this slave to the Jenkins master using the **Nodes** menu, giving it the label `docker`.
- 

# Phase 2: Kubernetes Node & Cluster Setup

This phase prepares the environment for our application deployment.

## 1. Launch Kubernetes Slave:

- Launch a new AWS EC2 instance.
- **Crucially, select instance type `t2.medium`** to meet the 2 CPU core requirement for Minikube.
- Name it `K8s-Slave`.

## 2. Install Prerequisites on K8s-Slave:

- SSH into the `K8s-Slave`.
- Install Java 17: `sudo yum install java-17-amazon-corretto -y`
- Install Docker (follow the same steps as for the Build Slave).
- Install `kubectl`:

Bash

```
curl -LO "https://dl.k8s.io/release/$(curl -L -s
https://dl.k8s.io/release/stable.txt)/bin/linux/amd64/kubectl"
chmod +x kubectl
sudo mv kubectl /usr/local/bin/
```

- **Install minikube:**

**Bash**

```
curl -Lo minikube
https://storage.googleapis.com/minikube/releases/latest/minikube-
linux-amd64
chmod +x minikube
sudo mv minikube /usr/local/bin/
```

### 3. Start the Kubernetes Cluster:

- Run the start command with the `--listen-address` flag. This is the key to making the cluster accessible from the EC2 instance's public IP.

**Bash**

```
minikube start --driver=docker --listen-address='0.0.0.0'
```

### 4. Connect K8s-Slave to Jenkins:

- In the Jenkins UI, go to **Manage Jenkins -> Nodes -> New Node**.
- Name it `K8s-Slave` and configure it as a permanent agent.
- Use the label `kubernetes`.
- Use the same SSH credentials and host key verification strategy as your other slaves.

---

## Phase 3: Jenkins Pipeline Configuration

This phase configures the Jenkins jobs to build and deploy the code.

### 1. Create a New GitHub Repository:

- Create a new repository (e.g., `jenkins-kubernetes-pipeline`).
- Add your `Dockerfile` and the new, improved `index.html` file to it.

### 2. Create the Build Job:

- In Jenkins, create a new Freestyle project named `K8s-Build-Job`.
- Restrict it to run on the `docker` label.
- Configure the **Source Code Management** section to point to your new GitHub repository.
- In **Build Environment**, bind your Docker Hub credentials to the `DOCKER_USER` and `DOCKER_PASS` variables.
- In the **Execute shell** build step, use this script:

## Bash

```
# Build the image
docker build -t my-k8s-app .

# Tag with build number and 'latest'
docker tag my-k8s-app $DOCKER_USER/my-k8s-app:$BUILD_NUMBER
docker tag my-k8s-app $DOCKER_USER/my-k8s-app:latest

# Login and push both tags
echo $DOCKER_PASS | docker login --username $DOCKER_USER --password-stdin
docker push $DOCKER_USER/my-k8s-app:$BUILD_NUMBER
docker push $DOCKER_USER/my-k8s-app:latest
```

### 3. Create the Kubernetes Deployment Job:

- Create a new Freestyle project named K8s-Deploy-Job.
- Restrict it to run on the kubernetes label.
- In the **Execute shell** build step, use this script. This script includes the critical rollout restart command to defeat caching issues.

## Bash

```
# Apply the deployment manifest
kubectl apply -f - <<EOF
apiVersion: apps/v1
kind: Deployment
metadata:
  name: my-web-app
spec:
  replicas: 1
  selector:
    matchLabels:
      app: my-web-app
  template:
    metadata:
      labels:
        app: my-web-app
    spec:
      containers:
        - name: my-web-container
          image: your-dockerhub-username/my-k8s-app:latest
          imagePullPolicy: Always
          ports:
            - containerPort: 80
EOF

# Force a rolling update to pull the new image
echo "--- Forcing a rolling restart of the deployment ---"
kubectl rollout restart deployment my-web-app
```

---

## Phase 4: Automation and Final Exposure

This final phase links everything together and makes the application public.

**1. Link the Jenkins Jobs:**

- Go to the configuration for the `K8s-Build-Job`.
- Add a **Post-build Action** of type **"Build other projects"**.
- Enter `K8s-Deploy-Job` in the "Projects to build" field.

**2. Set Up the GitHub Webhook:**

- In the `K8s-Build-Job` configuration, go to **Build Triggers** and check **"GitHub hook trigger for GITScm polling"**.
- In your GitHub repository settings, go to **Webhooks** and add a new webhook.
- The **Payload URL** must be  
`http://<Public_IP_of_Jenkins_Master>:8080/github-webhook/`.

**3. Expose the Application (Reliable Method):**

- SSH into your `K8s-Slave`.
- **Open a new, separate terminal window** and run the `kubectl port-forward` command. This must be left running.

Bash

```
kubectl port-forward deployment/my-web-app 8080:80 --  
address='0.0.0.0'
```

- In your **AWS Security Group** for the `K8s-Slave`, open port 8080 to Anywhere-IPv4.

**4. Access Your Application:**

- Open your browser and navigate to `http://<Public_IP_of_K8s_Slave>:8080`.